

```
print('hi')  
hi
```

## Question 1 of 4

### Basic Questions for class (2): Logistic Regression

Setting this up as a logistic regression  $y = \sigma(\beta_1 x$

- $\beta_0) y = \sigma(\beta_1 x + \beta_0)$  (where  $\sigma(x) = \frac{e^x}{1 + e^x}$ )
- $e^x \sigma(x) = 1 + e^x$

$e^x$

) problem:

What is the meaning of the input variable  $x$  here? What are the two parameters here? Why are they (and not the  $x$  values) parameters? What is the meaning of the output variable  $y$  here?

$x$  represents the input feature or independent variable used to predict the outcome. It's the observed data you have for each instance.

The two parameters are  $\beta_0$  (the intercept) and  $\beta_1$  (the coefficient for  $x$ ):  $\beta_0$ : This is the intercept. It represents the odds of the outcome when  $x=0$ .  $\beta_1$ : This is the weight or coefficient for the input variable  $x$ . It measures the change in the odds of the outcome for a one-unit change in  $x$ . They are parameters because they are the values the model learns from the data to best fit the relationship between  $x$  and  $y$ . The  $x$  values are not parameters; they are the input data you provide to the model.

$y$  is the output or dependent variable that the model aims to predict. In logistic regression,  $y$  is usually binary (0 or 1), representing the two possible classes or outcomes.

## Question 2 of 4: Likelihood

Remember the ( $y$ ) value we get from logistic regression really is a conditional probability. Specifically, it is

$$P(\text{True} \mid x) = \sigma(\beta_1 x + \beta_0)$$

From this, we can also deduce:

$$P(\text{False} \mid x) = 1 - \sigma(\beta_1 x + \beta_0)$$

Time for an analogy:

If a coin has a probability ( $p$ ) of landing heads up (and so,  $(1 - p)$  of landing tails up), and you observe heads, tails, heads, heads, tails, the likelihood is:

$$L(p) = p \cdot (1-p) \cdot p \cdot p \cdot (1-p)$$

And so the log likelihood is:

$$\ln L(p) = \ln(p) + \ln(1-p) + \ln(p) + \ln(p) + \ln(1-p)$$

**Question:** Fill in both expressions.

By analogy...the likelihood for our 100,000 observations ( $x_1, \dots, x_{100,000}$ ) is:

$$L(\beta_0, \beta_1) = \prod_{i=1}^{100,000} [\sigma(\beta_1 x_i + \beta_0)]^{y_i} [1 - \sigma(\beta_1 x_i + \beta_0)]^{1-y_i}$$

and so our log likelihood is:

$$\ln L(\beta_0, \beta_1) = \sum_{i=1}^{100,000} [y_i \ln \sigma(\beta_1 x_i + \beta_0) + (1 - y_i) \ln (1 - \sigma(\beta_1 x_i + \beta_0))]$$

## Core Question (3): Likelihood Calculation

Here's some code that "generates" some data, based on the probability distribution assumed by the logistic regression model. Specifically, the true values are  $\beta_0 = -5, \beta_1 = 3$  (but pretend you don't know that)

```
import numpy as np
from sklearn import metrics
import matplotlib.pyplot as plt

NUM_DATAPOINTS = 1000

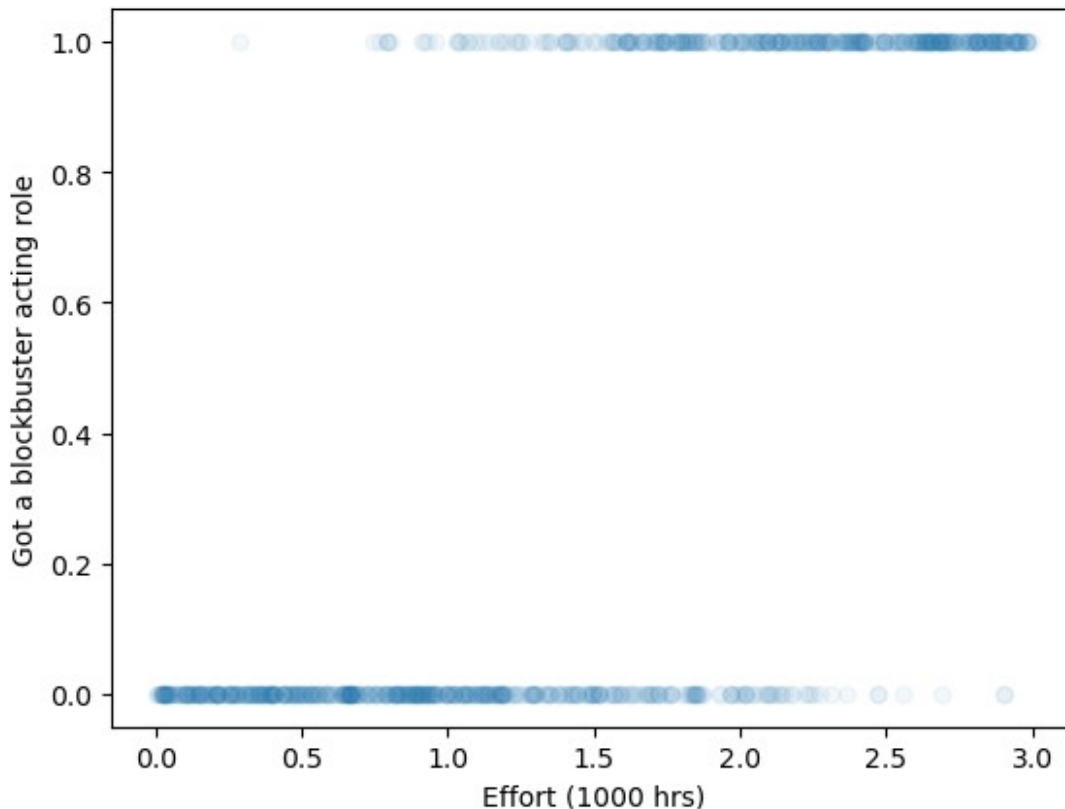
def sigmoid(x):
    return np.exp(x) / (1 + np.exp(x))

def generate_data(x):
    a = 3 #beta_1
    b = -5 #beta_0
    prob = sigmoid(a*x+b)
    return np.random.uniform() < prob

x_data = np.random.uniform(high=3, size=NUM_DATAPOINTS)
y_data = [generate_data(x) for x in x_data]
```

Note that this simulation produces a dataset where, we go from a low probability of success, to a high probability of success at about ~1200 hours, but also a fair amount of variance around that mark:

```
plt.scatter(x_data, y_data, alpha=0.05)
plt.xlabel("Effort (1000 hrs)")
plt.ylabel("Got a blockbuster acting role")
plt.show()
```



Now, suppose your friend Aishwarya (who doesn't know the real values of  $\beta_0$  and  $\beta_1$ ) claims the values that best "fit" this data are  $\beta_0 = 0.5$  and  $\beta_1 = 0.5$ . TA Irfan claims "well, actually..." the values are  $\beta_0 = 0.0$  and  $\beta_1 = 0.0$ . Both are wrong, but who is more wrong?

Computing the log-likelihood helps us answer this question: higher is better! Fill out the function below, and comment: which estimate is better?

```
# Fix this function
# You might wish to use the sigmoid(x) function defined above!
def log_likelihood_per_point(beta_0, beta_1, x, y):
    prob = sigmoid(beta_0 + beta_1 * x)
    if y is True:
        return np.log(prob)
    else:
        return np.log(1 - prob)

# This one's complete!
def total_log_likelihood(beta_0, beta_1, xs, ys):
    total_ll = 0

    for (x, y) in zip(xs, ys):
        total_ll += log_likelihood_per_point(beta_0, beta_1, x, y)
```

```

    return total_ll

# Be careful, which is the *higher* value when two values are
negative?
print("Aishwarya's estimate: ", total_log_likelihood( 3, -3, x_data,
y_data))
print("TA Irhum's estimate: ", total_log_likelihood( 4, -1, x_data,
y_data))

```

```

Aishwarya's estimate: -722.614375051468
TA Irhum's estimate: -2643.5425603316207

```

Aishwarya's estimate has a higher (less negative) log-likelihood than TA Irhum. Therefore, Aishwarya is less wrong, and her estimate better fits the data.

```

pip install jax

```

```

Collecting jax
  Using cached jax-0.4.33-py3-none-any.whl.metadata (22 kB)
Collecting jaxlib<=0.4.33,>=0.4.33 (from jax)
  Using cached jaxlib-0.4.33-cp312-cp312-
macosx_11_0_arm64.whl.metadata (983 bytes)
Collecting ml-dtypes>=0.2.0 (from jax)
  Using cached ml_dtypes-0.5.0-cp312-cp312-
macosx_10_9_universal2.whl.metadata (21 kB)
Requirement already satisfied: numpy>=1.24 in
/opt/anaconda3/lib/python3.12/site-packages (from jax) (1.26.4)
Collecting opt-einsum (from jax)
  Using cached opt_einsum-3.3.0-py3-none-any.whl.metadata (6.5 kB)
Requirement already satisfied: scipy>=1.10 in
/opt/anaconda3/lib/python3.12/site-packages (from jax) (1.13.1)
Using cached jax-0.4.33-py3-none-any.whl (2.1 MB)
Downloading jaxlib-0.4.33-cp312-cp312-macosx_11_0_arm64.whl (66.1 MB)
----- 66.1/66.1 MB 13.4 MB/s eta
0:00:0000:0100:01
l_dtypes-0.5.0-cp312-cp312-macosx_10_9_universal2.whl (750 kB)
----- 750.2/750.2 kB 17.0 MB/s eta
0:00:00a 0:00:01
-3.3.0-py3-none-any.whl (65 kB)
----- 65.5/65.5 kB 4.6 MB/s eta
0:00:00
, ml-dtypes, jaxlib, jax
Successfully installed jax-0.4.33 jaxlib-0.4.33 ml-dtypes-0.5.0 opt-
einsum-3.3.0
Note: you may need to restart the kernel to use updated packages.

import jax.numpy as jnp
from jax import grad
import numpy as np

```

```

# Sigmoid function using JAX
def sigmoid(x):
    return jnp.exp(x) / (1 + jnp.exp(x))

# Convert y_data from True/False to 0/1
y_data = np.array([int(y) for y in y_data])

# Function to compute the total log likelihood in JAX
def log_likelihood(beta_0, beta_1, x_data, y_data):
    linear_term = beta_1 * x_data + beta_0
    probs = sigmoid(linear_term)

    # Log likelihood: for y = 1, it's log(prob), for y = 0, it's log(1 - prob)
    return jnp.sum(y_data * jnp.log(probs) + (1 - y_data) * jnp.log(1 - probs))

# Compute gradient of the log likelihood w.r.t. beta_0 and beta_1
log_likelihood_grad = grad(log_likelihood, argnums=(0, 1))

# Aishwarya's initial estimate (for example, let's use these values)
initial_beta_0 = 3.0
initial_beta_1 = -3.0

# Compute the gradients at Aishwarya's initial estimate
grad_beta_0, grad_beta_1 = log_likelihood_grad(initial_beta_0,
initial_beta_1, x_data, y_data)

# Print the computed gradients
print(f"Gradient w.r.t. beta_0: {grad_beta_0}")
print(f"Gradient w.r.t. beta_1: {grad_beta_1}")

Gradient w.r.t. beta_0: 74.47683715820312
Gradient w.r.t. beta_1: 699.877197265625

```

## Question 3 of 4

## #MLDevelopment

Link your study Notion here and briefly describe the students you worked with and which materials you found particularly useful.

<https://www.notion.so/b96dd2ab4cd245748285211b477ebbb3?v=d31b2c53bf9a413583a1d286eeb94b0b&pvs=4>

I have worked with Hayoung to make this study Notion. <https://rpsychologist.com/likelihood/>  
This visualization of MLE calculation helped me have a deeper intuition about the whole process.